
Text-Aware Predictive Monitoring of Business Processes with LSTM Neural Networks

Master's Thesis

Author : **David Benedikt Georgi**

Supervisors : Dr.-Ing. M. Seran Uysal
M. Sc. Marco Pegoraro

Examiners : Prof. Dr. ir. Wil van der Aalst
Prof. Dr.-Ing. Ulrik Schroeder

Submission date : YYYY-MM-DD

This work is submitted to the institute

i9 Process and Data Science (PADS) Chair, RWTH Aachen University

Abstract

Contents

Abstract	ii
1 Introduction	1
1.1 Motivation	1
1.2 Problem Statement	2
1.3 Research Goals	2
1.4 Contribution	2
1.5 Thesis Structure	2
2 Preliminaries	3
2.1 Processes and Process Mining	3
2.2 Basic Notations and Sequences	4
2.3 Events, Traces, Event Logs	5
2.4 Text Mining	5
2.5 Supervised Learning	6
2.6 Long Short-Term Memory Networks	6
3 Related Work	9
4 Text-Aware Process Prediction	12
4.1 Overview	12
4.2 Event Encoding	14
4.3 Capturing Temporal Dependencies	14
4.4 Text Vectorization	15
4.4.1 Text Normalization	16
4.4.2 Bag of Words	16

4.4.3	Bag of N-Grams	17
4.4.4	Paragraph Vector	17
4.4.5	Latent Dirichlet Allocation	19
4.5	Network Architecture and Training	19
4.6	Predictive Business Processing Monitoring	20
5	Implementation	21
6	Evaluation	22
6.1	Datasets and	22
6.2	Next Event Prediction	22
6.3	Remaining Time Prediction	22
6.4	Outcome Prediction	22
7	Conclusion	23
	Bibliography	29
	Acknowledgement	30

Chapter 1

Introduction

1.1 Motivation

The rapid growth of data generated by large-scale information systems leads to new opportunities for society and businesses. In order to turn the massive amount of data into value, efficient techniques are needed, that are able to extract useful information. The generated value comes in different forms, for example visualizations, models, aggregated data or predictions can be of interest.

A remarkable subset of this data is described as *event data*, which is generated by *process-aware information systems* in order to manage, execute and monitor business processes [1]. With the non stopping rise of digitization of business processes more and more event data becomes utilizable, thus the potential value of this data is exploding.

The scientific engagement aiming to discover, analyze and improve real processes using event data led to *process mining*. Process mining bridges the gap between the data-driven characteristic of data science and the process-centric view of process science [2]. The ongoing success of process mining in research has been transferred to businesses, that successfully offer or utilize this technology. Celonis, which is often considered as one of the biggest commercial providers of process mining, has been valued 2.5 billion dollar only 9 years after the company was founded [3].

Modern process mining software tends to focus on continuous analysis rather than the traditional offline and project-based approaches. These business process monitoring systems are a key success factor for many organizations, since they allow to understand and supervise all connected processes of a company in real-time as the data is flowing. The core idea of this approach is to automate process mining and keep a persistent connecting between the business process data and the analytical capabilities.

However, traditional process mining tends to be backward-looking [4], i.e. it rather focuses on answering the question "What did happen?", rather than "What will happen?" or even "What should be done?". Therefore, new techniques are required to add the forward-perspective to process mining software.

1.2 Problem Statement

Businesses can develop a competitive advantage, if their process mining software has predictive capabilities, that allow to predict the future of a running process instance. For example, if it is known beforehand that a running process instance will probably exceed its deadline, measures can be initiated before damage occurs. Furthermore, information about the next event or the future path of process instance can be of interest. In some scenarios, processes instances have an outcome like success/failure or accepted/declined that can be predicted.

Precisely, given an event log with past executions of a process and a running (i.e. not completed) process instance, we would like to answer the following questions:

- What will happen next?
- When will it happen?
- What is the most likely future path of the instance?
- When will the instance finish?
- What is the outcome of the instance?

1.3 Research Goals

This thesis aims to improve current state-of-art approaches for process prediction to improve the capabilities of process monitoring software. The main research goal is to design, implement and evaluate a predictive model for event data that is able to take advantage of additional attribute and textual data associated with each event. Since most current approach are not able to handle textual data, we would like to know to which extend textual data can improve the quality of process prediction. Furthermore, we want to evaluate different design choices for a text-aware process prediction model and point out potential trade-offs.

1.4 Contribution

1.5 Thesis Structure

This thesis is structured in seven chapters. In Chapter 2 the notations, definitions and concepts used in this contribution are introduced. Chapter 3 summarizes relevant scientific contributions which focus on the problem of prediction in process mining to give an overview of already available methods and their capabilities. A novel text-aware process prediction model is presented in Chapter 4. Moreover, the details regarding the implementation of the model are given in Chapter 5. In Chapter 6 the performance of the model is evaluated and compared to current state-of-the-art prediction methods. Finally, in Chapter 7 a conclusion about all findings and an outlook towards future potential research on process prediction is given.

Chapter 2

Preliminaries

In this chapter the basic concepts of process mining, text mining, supervised learning, long short-term memory networks are presented. Furthermore, formal definitions and notations are introduced.

2.1 Processes and Process Mining

A *business process* is a collection of activities that are performed in a specific order to archive a goal [5]. A single execution of a process is a *case* or *process instance*, which is identified by a case ID. Each performed activity belongs to specific case and is completed at a certain time. For example, a case can be a patient in a hospital, a customer journey or an online order. The time on which an activity for a certain case is performed is specified by a timestamp. The trinity of case, activity and timestamp is called *event*. An event can have more attributes, for example resource, costs or transactional information.

If the execution of a business process is logged by an information system, the resulting event data is called *event log*. Depending on the format of the event log, it can also contain additional data on case level. Typical formats for event logs, are comma-separated values (CSV) and eXtensible Event Stream (XES) [6], which can be extracted from databases. A table-based representation of an artificial event log about patient treatment in a hospital can be seen in Table 2.1.

Process mining is the discipline that covers all approaches aiming to generate value out of event data. As an umbrella term, process mining includes or utilizes concepts of business process management, data mining, business process intelligence, big data, workflow management, business process monitoring [2] as well as machine learning [7].

Traditionally, process mining is divided into a set of subdisciplines mainly process discovery, conformance checking, process enhancement and process analytics [8]. Process discovery aims to generate process models out of event data in order to understand a process and enable further analysis. Conformance checking is about comparing the intended and observed behavior of a process. On top of these diagnostic approaches, process enhancement deals with the improvement of processes regarding compliance, performance or complexity.

ID	Activity	Timestamp	Resource	Cost	Comment
0	Register patient	01.02.2020:14.12	SYSTEM	0	-
	Consultation	01.02.2020:14.34	John Brown, MD	24.32	The patient reports persistent nausea.
	Blood test	01.02.2020:15.12	Kim Smith	14.23	Tests: Complete blood count
	Evaluate test result	01.02.2020:16.35	John Brown, MD	38.67	No abnormalities in the complete blood count.
	Release patient	01.02.2020:17.24	SYSTEM	0	-
1	Register patient	02.02.2020:08.20	SYSTEM	0	-
	Consultation	02.02.2020:14.12	Jana Simpson, MD	24.32	Noticeable tachycardia. No chronic diseases are known.
	MRI	02.02.2020:14.12	Sara Taylor, MD	352.87	-
	Release patient	02.02.2020:14.12	SYSTEM	0	-
2	Register patient	02.02.2020:09.08	SYSTEM	0	-
	Consultation	02.02.2020:09.14	Jana Simpson, MD	24.32	The patient has severe leg injuries due to a motorcycle accident.
	Patient hospitalized	02.02.2020:09.20	Mike Johnson	130.37	-
...

Table 2.1: Artificial event log of patient treatment in a hospital

Finally, process analytics focuses on metric and performance evaluation of processes. Similar to conformance checking, this term is closely related to business process monitoring, a rising subfield enabling the analysis of running business processes in real-time. Driven by the fast and ongoing development of quantitative prediction methods in data science and machine learning, also prediction-based methods have been applied to event data. These methods add the forward perspective to business process monitoring and deal with forecasting the future of a running process instance, which is also the main focus of this thesis.

2.2 Basic Notations and Sequences

The set \mathbb{N} denotes the set of all natural numbers $\{1, 2, 3, \dots\}$ and $\mathbb{N}_0 = \mathbb{N} \cup \{0\}$ denotes the set of natural numbers including 0. The set of natural numbers up to n is noted as $[n] = \{1, 2, \dots, n\} \subset \mathbb{N}$ with $[0] = \emptyset$.

Definition 2.1 (Sequence). A *sequence* of length $n \in \mathbb{N}_0$ over a set A is an ordered collection of elements defined by function $\sigma: [n] \rightarrow A$, which assigns each index an element of A . A sequence of length n is represented explicitly as $\sigma = \langle a_1, a_2, \dots, a_n \rangle$ with $a_i \in A$ for $1 \leq i \leq n$. In addition, $\langle \rangle$ is the empty sequence of length 0.

Given a set A , A^n describes the set of all sequences $\langle a_1, a_2, \dots, a_n \rangle$ over A of length n . The set A^0 is defined as $\{\langle \rangle\}$, the set that only contains the empty sequence. The set of all possible sequences over A is given with $A^* = \bigcup_{i \in \mathbb{N}_0} A^i$.

Given sequences σ_1 and σ_2 , the concatenation of both sequences is denoted by $\sigma_1 \cdot \sigma_2$. Moreover, the i -th element of a sequence $\sigma = \langle a_1, a_2, \dots, a_n \rangle$ is accessed using $\sigma(i) = a_i$ for $1 \leq i \leq n$. The length of a sequence is denoted by $|\sigma|$. For a sequence $\sigma = \langle a_1, a_2, \dots, a_n \rangle$, the function $hd^k(\sigma) = \langle a_1, a_2, \dots, a_k \rangle$ gives the prefix of length k of σ and $tl^k(\sigma) = \langle a_{n-k+1}, a_{n-k+2}, \dots, a_n \rangle$ the suffix of length k for $1 \leq k \leq n$.

A function $f: A \rightarrow B$ can be lifted element-wise to sequences over A , precisely:

$$f(\sigma) = \begin{cases} \langle \rangle & \text{if } \sigma = \langle \rangle \\ \langle f(a_1), f(a_2), \dots, f(a_n) \rangle & \text{else} \end{cases}$$

2.3 Events, Traces, Event Logs

Definition 2.2 (Event). An *event* is defined by tuple $e = (a, c, t, d_1, \dots, d_m) \in \mathcal{C} \times \mathcal{A} \times \mathcal{T} \times \mathcal{D}_1 \times \dots \times \mathcal{D}_m = \mathcal{E}$ where $c \in \mathcal{C}$ is the case ID, $a \in \mathcal{A}$ is the executed activity and $t \in \mathcal{T}$ is the timestamp of the event. Furthermore, each event contains a fixed number $m \in \mathbb{N}_0$ of additional attributes $d_1 \dots d_m$ in their corresponding domains $\mathcal{D}_1, \dots, \mathcal{D}_m$. In case that no additional attribute data is given ($m = 0$) the event space \mathcal{E} (set of all possible events) is reduced to $\mathcal{C} \times \mathcal{A} \times \mathcal{T}$.

Each attribute $d \in \mathcal{D}$ of an event (including activity, timestamp and case ID) can be accessed by a projection function $\pi_D: \mathcal{E} \rightarrow \mathcal{D}$. For example, the activity a of an event e is retrieved by $\pi_{\mathcal{A}}(e) = a$.

Throughout this thesis, $\mathcal{C} = \mathbb{N}_0$, $|\mathcal{A}| < \infty$ and $\mathcal{T} = \mathbb{R}$ is assumed, where $t \in \mathcal{T}$ is given in Unix time, precisely the number of seconds since 00:00:00 UTC on 1 January 1970 minus the applied leap seconds. Each additional attribute is assumed to be numerical, categorical or textual, i.e. $\mathcal{D}_i = \mathbb{R}$, $|\mathcal{D}_i| < \infty$ or $\mathcal{D}_i = \Sigma^*$ for $1 \leq i \leq m$ and some fixed Alphabet Σ .

Definition 2.3 (Trace). A *trace* is a finite and non-empty sequence of events $\sigma = \langle e_1, e_2, \dots \rangle \in \mathcal{E}^*$ with increasing timestamps, i.e. $\pi_{\mathcal{T}}(e_i) < \pi_{\mathcal{T}}(e_j)$ for $1 \leq i < j \leq |\sigma|$.

By lifting the projection functions to sequences a trace can be transformed into a sequence of attributes by applying the projection function to the trace. For example, $\pi_{\mathcal{A}}(\sigma)$ gives the sequence of the activities of the events in σ .

Definition 2.4 (Event log). An *event log* $\mathbb{L} = \{\sigma_1, \sigma_2, \dots, \sigma_k\}$ is a set of traces, where each event of a trace is unique in the log and all events of a trace share a case IDs, which is unique per trace.

2.4 Text Mining

Text mining describes all techniques to generate value out of unstructured or semi-structured textual data. It combines concepts of natural language processing, machine learning and data mining [9]. The base object in text mining is a *document* containing textual data. The text can be completely unstructured, i.e. it does not conform to a pre-defined data model, or semi-structured, like in an e-mail, where text information is assigned to sender, subject, message etc. A collection of documents is called *text corpus*, which forms the basis for many text mining techniques.

In order to derive a mathematical representation of the text data that can be interpreted by a computer, a text model has to be build using the text corpus. Popular text models are Bag-of-words, Bag-of-n-gram, Paragraph vector (a.k.a. Doc2Vec) [10] and Latent Dirichlet Allocation [11]. Most models require a text normalization step, where the text is cleaned from linguistic variation as well as meaningless words and symbols [12].

2.5 Supervised Learning

In *supervised learning* an unknown function is learned (i.e. approximated) from a set of example input-output pairs [13]. In contrast, *unsupervised learning* does not require examples pattern and is about finding pattern in the data. An input instance is usually described by a set of feature variables X and the output is defined by a target variable y . If the target variable y is continuous, we refer to this as a regression problem, if however it is discrete variable with a finite range of values, the learning problem is called classification problem. Given a *training set* of input-output pairs $\{(X_1, y_1), (X_2, y_2), \dots, (X_m, y_m)\}$, that were generated from an unknown function $y = f(X)$, the goal is to approximate a hypothesis function $h(X)$, which is close to $f(X)$, i.e. $h(X) \approx f(X)$.

The challenge in supervised learning is to generalize from the training set of input-output pairs in such a way, that the learned hypothesis function $h(x)$ can also successfully predict the target variable for unseen problem instances. In order to evaluate a hypothesis, the function is tested on a separate *test set* of input-output pairs, which has not been used for the construction of $h(X)$.

A hypothesis generalizes well, if its prediction performance is high on the training set as well as on test set. However, if the prediction performance is high on the training set, but not reliable on unseen data, the hypothesis *overfits* the training data. In this case, the model complexity, i.e. the number of parameters is higher than justified by the true function. In contrast, if the model is too simple to fit any data from training set, the hypothesis is *underfitting*.

In many real-world applications, the true function $f(X)$ is stochastic, i.e. we need to estimate a conditional probability function $P(Y|X)$ (classification problem) or a conditional expectation $E(Y|X)$ (regression problem) for prediction. Therefore, the prediction accuracy is always limited by the variation of the true distribution.

2.6 Long Short-Term Memory Networks

Long short-term memory (LSMT) is an advanced recurrent neural network architecture for sequential data originally presented by Hochreiter and Schmidhuber in 1997 [14]. This approach addresses the well-known vanishing and exploding gradient problem [15] of traditional recurrent neural networks by introducing more complex LSTM cells as hidden units. The proposed architecture has been improved several times [16] [17] and considered as one of the most successful recurrent neural network models. Although LSTM networks have been available for a long time, the breakthrough of this technology is dated around 2016 after many success stories of LSTM in combination with large data sets and GPU hardware have been reported for sequence to sequence tasks like text translation [18].

Gated recurrent units (GRU) [19] are the competing gating mechanism by Cho et al. that have fewer parameters and perform similar to LSTM. However, more recent studies show, that LSTM outperforms GRU consistently in neural machine translation tasks [20].

A simple feedforward neural networks consists of an input layer, arbitrarily many hidden layers and an output layer, where each layer consists of neurons that compute and output the weighted sum of the cells of the previous layer that has been passed to a non-linear activation function [21]. These networks can learn and compute complex functions in supervised learning settings, where input and output pattern are provided. The network computes a loss function for each training pattern and adjusts its weights with gradient descents using a back-propagation algorithm in order to minimize the loss function [22].

Recurrent neural networks extend traditional feed forward networks with backfeeding connections between hidden layers. This enables the network to keep a state across inputs and allows the neural network to process arbitrarily long sequences of input data while learning temporal dependencies.

In LSTM networks the layers are replaced by more complex LSTM modules, where each module contains four different sublayers. The module uses as input the state \mathbf{c}_{t-1} and the hidden output \mathbf{h}_{t-1} of the module in the previous time step as well as the output of the previous layer \mathbf{x}_t to compute a new cell state \mathbf{c}_t and a (hidden) output \mathbf{h}_t .




Figure 2.1: An LSTM module with four sublayers that manipulate the cell state and compute the module's output. Graphic adapted from [23].

The input vector \mathbf{x}_t is concatenated with the previous hidden output \mathbf{h}_{t-1} and fed to four neural network layers, which are designed to decide what part of the cell state will remain (forget gate f_t), how it is updated (update gate i_t and $\tilde{\mathbf{c}}_t$) and what the output of the layer will be (output gate o_t leading to \mathbf{h}_t). The sublayer apply $\text{sigmoid}(x) = \frac{1}{1+\exp(-x)}$

or $\tanh(x) = \frac{\exp(x) - \exp(-x)}{\exp(x) + \exp(-x)}$ activation functions elementwise, leading to the following equations:

$$\mathbf{f}_t = \text{sigmoid}(\mathbf{W}_f \cdot (\mathbf{h}_{t-1}, \mathbf{x}_t) + \mathbf{b}_f)$$

$$\mathbf{i}_t = \text{sigmoid}(\mathbf{W}_i \cdot (\mathbf{h}_{t-1}, \mathbf{x}_t) + \mathbf{b}_i)$$

$$\bar{\mathbf{c}}_t = \tanh(\mathbf{W}_c \cdot (\mathbf{h}_{t-1}, \mathbf{x}_t) + \mathbf{b}_c)$$

$$\mathbf{o}_t = \text{sigmoid}(\mathbf{W}_o \cdot (\mathbf{h}_{t-1}, \mathbf{x}_t) + \mathbf{b}_o)$$

\mathbf{W}_f , \mathbf{W}_i , \mathbf{W}_c and \mathbf{W}_o are the sublayer's learned weights and \mathbf{b}_f , \mathbf{b}_i , \mathbf{b}_c and \mathbf{b}_o are the corresponding biases.

The new cell state \mathbf{c}_t is then a combination of the old cell state \mathbf{c}_{t-1} and the result of the update gate $\bar{\mathbf{c}}_t$, where the layer computations \mathbf{f}_t and \mathbf{i}_t determine the proportions by a pointwise multiplication (\odot) with the cell states.

$$\mathbf{c}_t = \mathbf{f}_t \odot \mathbf{c}_{t-1} + \mathbf{i}_t \odot \bar{\mathbf{c}}_t$$

The result of the output gate \mathbf{o}_t is pointwise multiplied with the tanh-activated new cell state to calculate the hidden output \mathbf{h}_t of the module.

$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{c}_t)$$

LSTM networks are able to backpropagate a more stable error with this gating mechanism, such that these networks are much more capable of learning complex functions for sequences compared to standard recurrent neural networks.

Chapter 3

Related Work

The prediction of the future of an process instance has been an important subfield in the process mining research, that aims to enhance process monitoring capabilities. Depending on the use case, for example predicting time-related attributes, the future path or the the outcome of a case can be of interest. Most approaches presented in the literature either use machine learning models or process models to construct a predictor that generalizes from a historical event log.

van Dongen et al. presented five different non-parametric regression predictors for forecasting the total cycle time of an unfinished case[24]. The estimates are based on activity occurrences, activity duration and attributes.

van der Aalst et al. proposed to build a transition system using a set, bag or sequence abstraction, which is annotated with time-related data in order to predict the remaining time of case [25]. The core idea of this approach is to replay unfinished cases on the learned transition system and compute the prediction using the annotated data.

Pandey et al. use a hidden markov model to predict the remaining time of a case using the activity and timestamp data of an event log [26].

Rogge-Solti and Weske showed how a stochastic Petri net can be used to predict the remaining of a process instance. The model naturally supports parallelism in business processes and considers future events which are expected to occur.

Ceci et al. presented an approach, where a sequence tree is learned in order relate a running traces to similar historical traces [28]. A decision tree is then used to predict the next activity and the remaining time of a case.

Teinemaa et al. applied text vectorization techniques like bag-of-n-grams (BoNG), Latent Dirichlet Allocation (LDA) and Paragraph Vectors (PV) to textual data of processes in order to predict a binary label describing the process outcome[29]. In this approach random forest and logistic regression classifiers for each prefix length of a trace are trained.

Most recently, several authors have applied recurrent neural networks in form of LSTM networks for process prediction. Evermann et al. encode events using an embedding matrix as it is known for word embeddings. The embedded events are then used as input for an LSTM network that predicts the next activity[30].

Tax et al. use an one-hot-encoding of the activity and the timestamp of an event to predict the activity and timestamp of the next event. This is done by using a two-layered LSTM network[31].

The work by Navarin et al. adopts the idea of using an LSTM network [31] and extends the encoding to additional data attributes associated with each event[32] to predict the remaining time of an case.

Polato et al. presented a set of approaches that use support vector regression for remaining time prediction[33]. In this work the authors implement different encoding for events including simple one-hot-encoding and a more advanced state based encoding using transition systems. Furthermore, they enhance the approach in [25] by taking additional data attributes into account.

Teinemaa et al. reported an in-depth review and benchmark of outcome-oriented predictive process monitoring approaches. The study showed that aggregated encoding like counting frequencies of activities as most reliable encoding for outcome-prediction [34].

Park and Song showed how LSTM-based predictions can be used to solve a resource allocation problem, leading to direct recommendations for process improvement [35].

A comparison of the process prediction methods is presented in Table 3.1.

Contribution	Year	Model(s)	Data-Aware	Text-Aware	Predictions
van Dongen et al. [24]	2008	Regression	✓	✗	Remaining time
van der Aalst et al. [25]	2011	Transition system	✗	✗	Remaining time
Pandey et al. [26]	2011	Hidden Markov	✗	✗	Remaining time
Rogge-Solti and Weske [27]	2013	Stochastic Petri Net	✗	✗	Remaining time
Ceci et al. [28]	2014	Sequence Tree Decision Tree	✓	✗	Next activity Remaining time
Teinemaa et al. [29]	2016	Random Forest Logistic regression	✓	✓	Case outcome
Evermann et al. [30]	2016	LSTM	✗	✗	Next activity
Tax et al. [31]	2017	LSTM	✗	✗	Next activity Future path Next event time Remaining time
Navarin et al. [32]	2017	LSTM	✓	✗	Remaining time
Polato et al. [33]	2018	Transition system SVR Naive Bayes	✓	✗	Next activity Future path Remaining time
Park and Song [35]	2019	LSTM	✓	✗	Next activity Next event time
This contribution	2020	LSTM	✓	✓	Next activity Future path Next event time Remaining time Case outcome

Table 3.1: Comparison of process prediction methods.

Chapter 4

Text-Aware Process Prediction

Text-aware process prediction aims to utilize unstructured text information in event data to improve predictions for unfinished cases. While many prediction methods have been applied to event data, almost none of them are able to handle textual data. A first approach has been presented in [29], where traces with text data are encoded as vectors and a random forest classifier is learned for each prefix length.

In this chapter a novel approach for text-aware process prediction is presented that considers control flow, additional numerical, categorical and textual data, captures temporal dependencies between events, seasonal variability and concept drifts using an event-wise encoding and a sequential LSTM prediction model.

4.1 Overview

The proposed framework consists of a preprocessing, encoding and prediction model component, which operate in an offline and online phase. In the offline phase a historical event log with completed traces of a business process is used to fit the encoding and prediction component. Given an event log $\mathbb{L} = \{\sigma_1, \dots, \sigma_l\}$ with historical traces, the set of all prefix traces $\mathbb{L}_{\text{prefix}} = \{hd^k(\sigma) \mid \sigma \in \mathbb{L}, 1 \leq k \leq |\sigma|\}$ is computed and encoded as a sequence of event vectors. The encoding component distinguishes between categorical or numerical data that can be encoded directly and textual data, which is encoded based on a textual model. Each encoded prefix sequence with its desired target values corresponds to one training example for an LSTM network.

The target values of a prefix sequence are the activity and timestamp (relative to case start) of the next event as well as the outcome and cycle time (time between first and last event). For completed cases the next activity is an artificial "<Process End>" activity with the same timestamp as the final event.

The total number of training examples that can be generated out of the log is $\sum_{\sigma \in \mathbb{L}} |\sigma|$, which is exactly the number of events in the log.




Figure 4.1: Framework

4.2 Event Encoding

In the offline training phase as well as during online prediction, traces are encoded as sequences of event vectors. Strictly speaking, a set of encoding functions $enc_k: \mathcal{E}^k \rightarrow (\mathbb{R}^n)^k$ are realized by the encoding component, that encodes (prefix-)traces of length k to vector sequences of the same size. Each event is encoded as a fixed-length vector using the activity, timestamp and additional categorical, numerical and textual data. Each encoded event vector is the concatenation of set of feature vectors, which are constructed from the event data.

The activity of an event is represented using *one-hot encoding*. Given the set of possible activities \mathcal{A} , an arbitrary but fixed ordering over is introduced with a bijective index function $index_{\mathcal{A}}: \mathcal{A} \rightarrow \{1, \dots, |\mathcal{A}|\}$. Using this function, the activity is encoded as a vector of size $|\mathcal{A}|$, where the component $index_{\mathcal{A}}(\pi_{\mathcal{A}}(e))$ has value 1 and all the other components have value 0. We write $\mathbb{1}$

The timestamp of an event is used to compute a set of time-related features, which is

$$y_a = \mathbb{1}(hd^1(tl^k(\pi_{\mathcal{A}}(\sigma))))$$

$$y_t = \mathbb{1}(hd^1(tl^k(\pi_{\mathcal{T}}(\sigma))))$$

$$y_{outcome} = \mathbb{1}(hd^1(tl^k(\pi_{\mathcal{T}}(\sigma))))$$

$$y_{cycle} = \pi_{\mathcal{T}}(\sigma(|\sigma|)) - \pi_{\mathcal{T}}(\sigma(1))$$

$$\mathbf{x} = (\mathbf{a}, \mathbf{t}, \mathbf{d}_1^{\text{num}}, \dots, \mathbf{d}_r^{\text{num}}, \mathbf{d}_1^{\text{cat}}, \dots, \mathbf{d}_s^{\text{cat}}, \mathbf{d}_1^{\text{text}}, \dots, \mathbf{d}_u^{\text{text}}) \quad (4.1)$$

Feature Vector	Construction	Dimension	Description
\mathbf{a}	$\mathbb{1}(\pi_{\mathcal{A}}(e))$	$ \mathcal{A} $	One-hot encoding of the activity.
\mathbf{t}	See Section 4.3	6	Time-based feature vector.
$\mathbf{d}_i^{\text{num}}$	$norm(\pi_{\mathcal{D}_i^{\text{num}}}(e))$	1	Normalized value of the i -th numerical attribute
$\mathbf{d}_i^{\text{cat}}$	$\mathbb{1}(\pi_{\mathcal{D}_i^{\text{cat}}}(e))$	$ \mathcal{D}_i $	One-hot encoding of the i -th categorical attribute.
$\mathbf{d}_i^{\text{text}}$	See Section 4.4	z_i	Fixed-length vectorization of the i -th text attribute.

Table 4.1: Feature vectors for event encoding.

$$|\mathbf{x}| = |\mathcal{A}| + r + \sum_{i=1}^s |\mathcal{D}_i| + \sum_{j=1}^u z_j + 6 \quad (4.2)$$

4.3 Capturing Temporal Dependencies

For timestamp prediction of the next and final event of running process instance a set of time-based features is computed from the timestamp data in the event log. Given an event

log \mathbb{L} an event e_i from a trace $\sigma = \langle e_1, \dots, e_n \rangle$ the following time features are computed for the encoding of e_i :

Feature	Description
$t_1 = \pi_{\mathcal{T}}(e_i) - \pi_{\mathcal{T}}(e_{i-1})$	Seconds since previous event
$t_2 = \pi_{\mathcal{T}}(e_i) - \pi_{\mathcal{T}}(e_1)$	Seconds since case start
$t_3 = \pi_{\mathcal{T}}(e_i) - \min\{\pi_{\mathcal{T}}(e_j) \mid e_j \in \sigma_k, \sigma_k \in \mathbb{L}\}$	Seconds since first recorded event
t_4	Seconds since midnight
t_5	Seconds since last Monday
t_6	Seconds since last January 1 00:00

Using the time features a set of time-dependent trends can be captured and utilized for prediction. The features t_1 and t_2 give information about the time between events and the time of the event in the case. Using t_3 the absolute time position of an event in the data can be determined. This is important to detect concept drift in the process [36]. The features t_4, t_5 and t_6 are used to capture daily, weekly or seasonal trends. For example, some activities might only be executed during office hours, before the weekend, during summer.

Each feature t_1, \dots, t_6 as well as all additional numerical attributes d_i are scaled to the interval $[0, 1]$ to improve learning efficiency using min-max normalizing. The scaling for a numerical feature x is realized with the transformation

$$\hat{x} = \frac{x - \min(x)}{\max(x) - \min(x)},$$

where $\min(x)$ is the lowest and $\max(x)$ is the highest value x can take. If the limits are not bounded conceptually, the lowest or highest value of x in the event log is used for scaling.

4.4 Text Vectorization

In order to prepare the textual data of the event log for a prediction model, the texts have to be encoded in a compact, finite and "useful" numerical vector representation using a text model. Useful in that context means, that texts with similar semantic meanings should also have similar representations. The vector representation of text data is an important aspect in *Natural Language Processing* (NLP). Extracting the meaning of textual information remains a challenge even for humans, because textual data is unstructured, language dependent and domain specific. Many words are ambiguous, for example the word "apple" might denote a fruit or a global technology company. In addition, grammatical variations and the importance of context in language makes extracting the semantic meaning even more difficult for computers.

In our setting, the text vectorization for textual attributes is realized in a two step procedure by text encoding component. First, all text data associated with the events in the corresponding textual attribute is collected in a so called *text corpus*. Each document in the text corpus is then preprocessed in order to filter out linguistic noise or useless information. This step is called *text normalization*. Finally, the text corpus is used to build up a *vocabulary* and a text vectorization technique is applied to encode the text

attribute into a fixed-length vector. As text vectorization techniques, the Bag of Words, Bag of N-Grams, Paragraph Vector and Latent Dirichlet Allocation are considered. The vocabulary of a text corpus is a set V of all relevant words that appear in the corpus and is usually indexed by an bijective index function $index_V: V \rightarrow \{1, 2, \dots, |V|\}$.

4.4.1 Text Normalization

In the text normalization step each document is transformed by a preprocessing pipeline which consists of the following four steps:

1. Letters are converted to lowercase
2. Document is tokenized (i.e. splitted) by word
3. Each word is lemmatized
4. Stop words are filtered out

In the tokenenization step a document is split up in a list of words. Each word is then lemmatized, i.e. it is converted to its canonical form. The idea is to unify words that have a very similar meaning and filter out grammatical variations. For example, the words "go", "going", "goes", "gone" and "went" are all transformed to the basic form "go".

Ultimately, all stop words are filtered out of each document. Stop words are words with low information value like "the", "a", "of" or "here". Stop word lists are language dependent and can be more or less aggressive at filtering. Usually they contains articles, auxiliary verbs, prepositions and generic verbs like "be" and "have". In addition, punctuation marks or numerical information are excluded.

Step	Transformation	Example Document
0	Original	"The patient has been diagnosed with high blood pressure."
1	Lowercase	"the patient has been diagnosed with high blood pressure."
2	Tokenization	["the", "patient", "has", "been", "diagnosed", "with", "high", "blood", "pressure", "."]
3	Lemmatization	["the", "patient", "have", "be", "diagnose", "with", "high", "blood", "pressure", "."]
4	Stop word filtering	["patient", "diagnose", "high", "blood", "pressure"]

Table 4.2: Preprocessing transformation of an example document containing a single sentence.

4.4.2 Bag of Words

The *Bag of Words* (BoW) Model is a simple text model, which represents documents based on the *term frequencies* of its words [37]. Given the learned vocabulary V a document is represented by a vector of size $|V|$, where the i -th component gives the number of occurrences of the word in the document indexed with i it.

Since this approach does not reflect the prior distributions of words in the corpus, i.e. how likely certain words appear in a document, the term frequencies are usually normalized by the so-called *inverse document frequency* (idf) of a word. The inverse document frequency

indicates the specificity of a word in the corpus and is computed by dividing the total number of documents by the number of documents that contain the specific word and scaling that value logarithmically. The resulting score is the tfidf score of a word in a document.

The Bag of Words model is easy to build and effective for certain applications, but limited in several ways. First, the model completely ignores the order of words, which is often crucial for understanding the semantic meaning. For example, the sentences "The patient's health state went from stable to critical." and "The patient's health state went from critical to stable." would result in the same vector representation, while the meaning is clearly inverted. Second, the vector representations are sparse and of high dimensionality since they depend on the size of the vocabulary. However, the dimension can be reduced by limiting the size of the vocabulary. For example, words with low tfidf scores can be excluded from the vocabulary.

4.4.3 Bag of N-Grams

The *Bag of N-Grams* model is a generalization of the Bag-of-Words model, which addresses the missing word order awareness of the latter. Instead of single words, the vocabulary consists of n -tuples of words, that appear consecutive in the corpus. The unigram model ($n = 1$) is equivalent to the BoW model. For the bigram model ($n = 2$), the vocabulary consists of pairs of words that appear next to each other in the documents. For example, for our preprocessed document ["patient", "diagnose", "high", "blood", "pressure"], the pairs ("patient", "diagnose"), ("diagnose", "high"), ("high", "blood") and ("blood", "pressure") are taken into the vocabulary. For $n > 2$ n -tuples are generated accordingly. The feature vector is constructed by computing the tfidf score for each vocabulary entry like in the BoW model.

Compared to the BoW model, n -grams also take the order of words into account, which is beneficial in many scenarios. However, the vocabulary size is usually even higher than in the BoW model. In order to generate more compact vectors, distributed text representations are needed for larger text corpora and vocabularies.

4.4.4 Paragraph Vector

The *Paragraph Vector* also known as *Doc2Vec*, originally presented by Le and Mikolov, is an unsupervised algorithm that learns distributed fixed length vector representations for documents of variable length [10]. The idea is inspired by the word embedding model presented by Bengio et al. [38], which can learn distributed fixed-length vector representations for words. In this model words are mapped to vectors, that are trained to predict words from its context, i.e. words that appear before or after the target word in the training documents. Several variants of this approach exist, notably the *Continuous Bag of Words* model, which ignores the order of the words in the context and the *Continuous Skip-gram* model, which predicts the skip-gram context for a word vector (also known as *Word2Vec*) [39].

The core idea of the Paragraph Vector model is to extend the model in [38] in a way, that an additional document or paragraph vector, that is unique per document is trained together with the word vectors. Fig. 4.2 shows the architecture of the distributed memory variant of the Paragraph Vector model (PV-DM). It is realized by a neural network, that takes

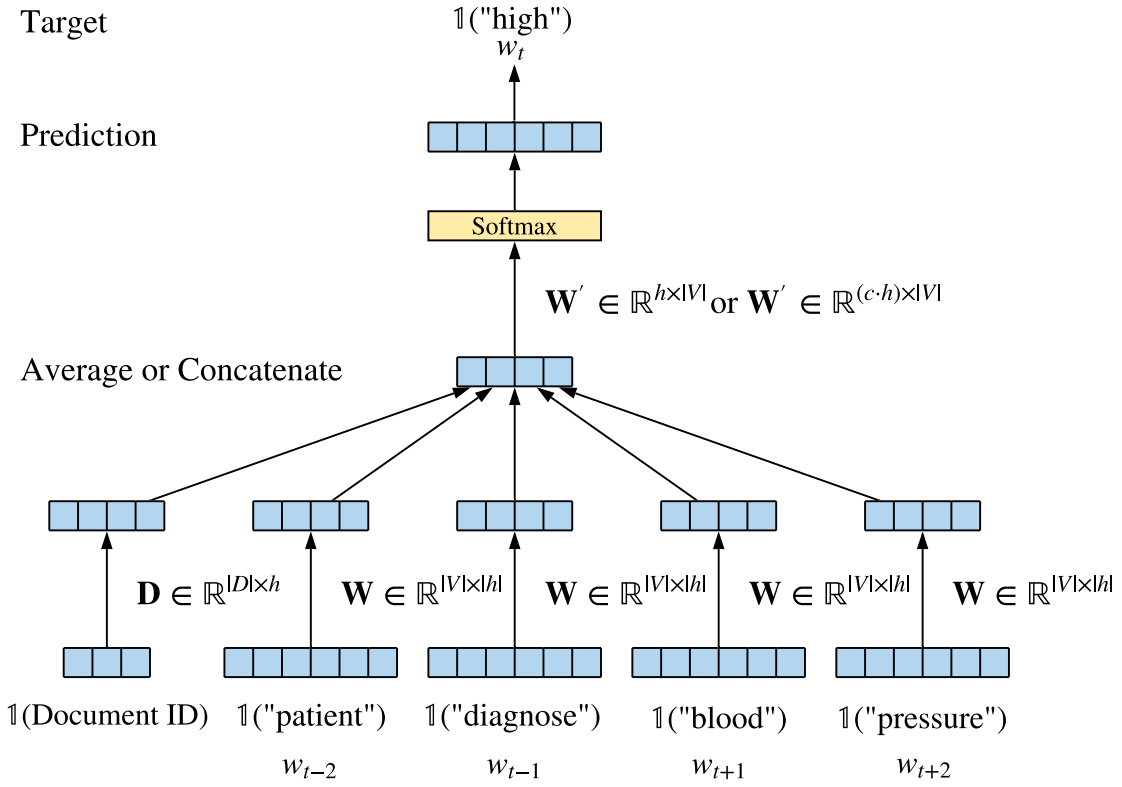


Figure 4.2: The distributed memory Paragraph Vector model (PV-DM) is designed to predict a word w_t from its context and derives fixed-length representation of documents and words via the learned matrices \mathbf{D} and \mathbf{W} .

one-hot encoded words and a one-hot encoded document IDs as input. These are mapped to vector representation via weight matrices \mathbf{D} and \mathbf{W} , which are learned during training with gradient descent. The distributed representations are then averaged or concatenated to a vector in order to predict the one-hot encoded target word using another mapping via \mathbf{W}' and a softmax activation function. The training set is constructed using a sliding window over every document, such that the input is the context of the target word and the document ID. After training, each column in \mathbf{D} represents the distributed encoding of the corresponding document.

The network is also able to learn a representation for new unseen documents by an inference step. In this phase, the word matrix \mathbf{W} and the prediction matrix \mathbf{W}' are fixed and only the document vector is trained. The Paragraph Vector model tends to perform better than non-distributed models, however since new documents are vectorized via inference, a bigger training corpus of documents is usually required.

4.4.5 Latent Dirichlet Allocation

4.5 Network Architecture and Training

The LSTM network is designed to be trained with all prediction targets (next activity, next event time, case outcome and case cycle time) at once, in order to benefit from correlations between these. The network consists of an input layer, an shared LSTM layer, an a specialized LSTM layer for each target, and an fully connected output layer for each target. Furthermore, layer normalization [40] is applied after each LSTM layer, which standardizes the hidden output in order to speed up the training convergence.

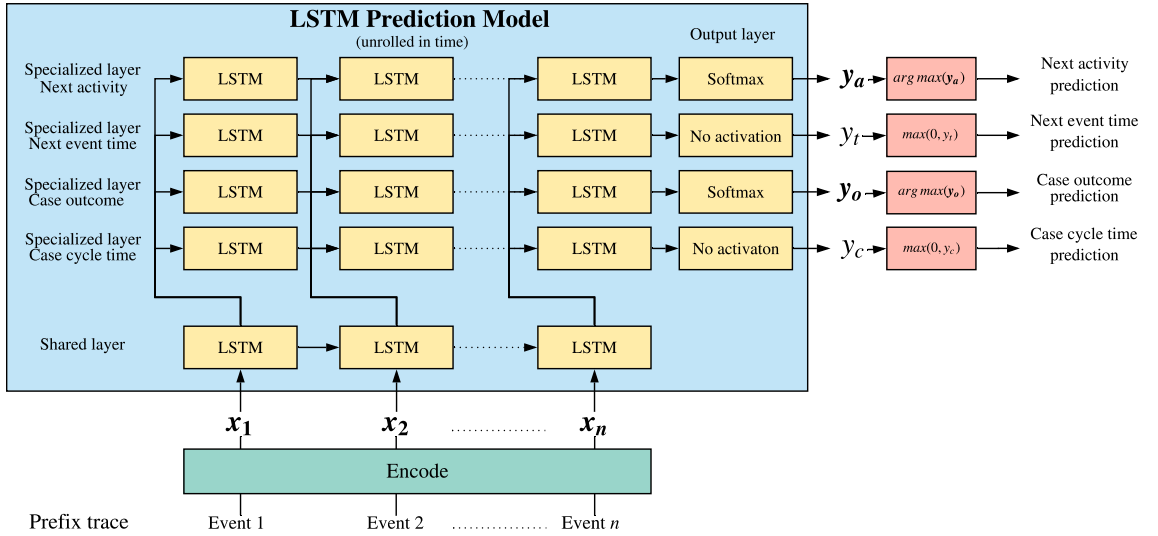


Figure 4.3: LSTM model to simultaneously predict the next activity (y_a), next event time (y_t), case outcome (y_o) and case cycle time (y_c) for an encoded prefix trace x_1, x_2, \dots, x_n .

The fully connected output layer uses a softmax activation function for the next activity and case outcome prediction to estimate the probability for each target value. The function normalizes a vector of real numbers into another vector of same dimension such that all

component are in the interval (0,1) and the sum of all component is equal to 1. Hence, the transformed vector can be interpreted as a probability distribution, while keeping the proportions of the original vector. The softmax function is described with

$$\text{Softmax}(\mathbf{x})_i = \frac{\exp(x_i)}{\sum_{j=1}^n \exp(x_j)} \text{ for } i = 1, \dots, n \text{ and } \mathbf{x} = (x_1, \dots, x_n) \in \mathbb{R}^n.$$

The training set of encoded prefix traces is represented by an 3-dimensional matrix of real values, where the three dimensions specify the prefix traces, the events per prefix trace and the features per event. Since the prefix traces have different length, shorter traces are pre-padded [41] with zeros vectors. Hence, a prefix trace of encoded events $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n$ is represented in the training set by a 2-dimensional matrix $(\mathbf{0}, \dots, \mathbf{0}, \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n)$, such that the zero vectors fill up shorter traces to the length of the longest trace in the training set. All prefix traces together form the 3-dimensional training matrix.

The training is realized by a backpropagation through time (BPTT) algorithm that updates the weights of the network using the update rules of the Adam optimizer with Nesterov momentum [42]. The loss for numerical prediction values \hat{y} and the true value y is the absolute error $\text{AE}(\hat{y}, y) = |\hat{y} - y|$, while the loss for categorical prediction values is computed using the categorical cross entropy error $\text{CE} = -\sum_{i=1}^k y_i \cdot \log \hat{y}_i$.

4.6 Predictive Business Processing Monitoring

During online business process monitoring, predictions are realized by a forward-pass of the encoded running cases through the LSTM model. The component with the highest value of the softmax outputs for the next activity (\mathbf{y}_a) and the case outcome (\mathbf{y}_o) indicates the categorical prediction. The output values for the next event time (y_t) and case duration (y_c) are clipped to 0 for negative outputs and the normalization is reverted, in order to compute the final prediction value.

Chapter 5

Implementation

Chapter 6

Evaluation

6.1 Datasets and

6.2 Next Event Prediction

6.3 Remaining Time Prediction

6.4 Outcome Prediction

Chapter 7

Conclusion

Bibliography

- [1] Wil M. P. van der Aalst. Process-aware information systems: Lessons to be learned from process mining. *Trans. Petri Nets Other Model. Concurr.*, 2:1–26, 2009. doi: 10.1007/978-3-642-00899-3_1. URL https://doi.org/10.1007/978-3-642-00899-3_1.
- [2] Wil M. P. van der Aalst. *Process Mining - Data Science in Action, Second Edition*. Springer, 2016. ISBN 978-3-662-49850-7. doi: 10.1007/978-3-662-49851-4. URL <https://doi.org/10.1007/978-3-662-49851-4>.
- [3] Ryan Browne. How three friends turned a college project into a \$2.5 billion software unicorn, 2019. URL <https://web.archive.org/web/20200125191917/https://www.cnbc.com/2019/11/21/celonis-raises-290m-series-c-funding-round-at-2point5b-valuation.html>. Accessed = 2020-04-20.
- [4] Wil M. P. van der Aalst. Process mining and simulation: a match made in heaven! In *Proceedings of the 50th Computer Simulation Conference, Summer-Sim 2018, Bordeaux, France, July 09-12, 2018*, pages 4:1–4:12. ACM, 2018. URL <http://dl.acm.org/citation.cfm?id=3275386>.
- [5] Wil M. P. van der Aalst, Arya Adriansyah, Ana Karla Alves de Medeiros, Franco Arcieri, Thomas Baier, Tobias Blickle, R. P. Jagadeesh Chandra Bose, Peter van den Brand, Ronald Brandtjen, Joos C. A. M. Buijs, Andrea Burattin, Josep Carmona, Malú Castellanos, Jan Claes, Jonathan Cook, Nicola Costantini, Francisco Curbera, Ernesto Damiani, Massimiliano de Leoni, Pavlos Delias, Boudewijn F. van Dongen, Marlon Dumas, Shahram Dustdar, Dirk Fahland, Diogo R. Ferreira, Walid Gaaloul, Frank van Geffen, Sukriti Goel, Christian W. Günther, Antonella Guzzo, Paul Harmon, Arthur H. M. ter Hofstede, John Hoogland, Jon Espen Ingvaldsen, Koki Kato, Rudolf Kuhn, Akhil Kumar, Marcello La Rosa, Fabrizio Maria Maggi, Donato Malerba, R. S. Mans, Alberto Manuel, Martin McCreesh, Paola Mello, Jan Mendling, Marco Montali, Hamid R. Motahari Nezhad, Michael zur Muehlen, Jorge Munoz-Gama, Luigi Pontieri, Joel Ribeiro, Anne Rozinat, Hugo Seguel Pérez, Ricardo Seguel Pérez, Marcos Sepúlveda, Jim Sinur, Pnina Soffer, Minseok Song, Alessandro Sperduti, Giovanni Stilo, Casper Stoel, Keith D. Swenson, Maurizio Talamo, Wei Tan, Chris Turner, Jan Vanthienen, George Varvaressos, Eric Verbeek, Marc Verdonk, Roberto Vigo, Jianmin Wang, Barbara Weber, Matthias Weidlich, Ton Weijters, Lijie Wen, Michael Westergaard, and Moe Thandar Wynn. Process mining manifesto. In Florian Daniel, Kamel Barkaoui, and Shahram Dustdar, editors, *Business Process Management Workshops - BPM 2011 International Workshops, Clermont-Ferrand*,

- France, August 29, 2011, *Revised Selected Papers, Part I*, volume 99 of *Lecture Notes in Business Information Processing*, pages 169–194. Springer, 2011. doi: 10.1007/978-3-642-28108-2_19. URL https://doi.org/10.1007/978-3-642-28108-2_19.
- [6] Eric Verbeek, Joos C. A. M. Buijs, Boudewijn F. van Dongen, and Wil M. P. van der Aalst. XES tools. In Pnina Soffer and Erik Proper, editors, *Proceedings of the CAiSE Forum 2010, Hammamet, Tunisia, June 9-11, 2010*, volume 592 of *CEUR Workshop Proceedings*. CEUR-WS.org, 2010. URL <http://ceur-ws.org/Vol-592/PaperDemo07.pdf>.
- [7] Fabian Veit, Jerome Geyer-Klingenberg, Julian Madrzak, Manuel Haug, and Jan Thomson. The proactive insights engine: Process mining meets machine learning and artificial intelligence. In Robert Clarisó, Henrik Leopold, Jan Mendling, Wil M. P. van der Aalst, Akhil Kumar, Brian T. Pentland, and Mathias Weske, editors, *Proceedings of the BPM Demo Track and BPM Dissertation Award co-located with 15th International Conference on Business Process Modeling (BPM 2017), Barcelona, Spain, September 13, 2017*, volume 1920 of *CEUR Workshop Proceedings*. CEUR-WS.org, 2017. URL http://ceur-ws.org/Vol-1920/BPM_2017_paper_192.pdf.
- [8] Maikel L. van Eck, Xixi Lu, Sander J. J. Leemans, and Wil M. P. van der Aalst. PM²: A process mining project methodology. In Jelena Zdravkovic, Marite Kirikova, and Paul Johannesson, editors, *Advanced Information Systems Engineering - 27th International Conference, CAiSE 2015, Stockholm, Sweden, June 8-12, 2015, Proceedings*, volume 9097 of *Lecture Notes in Computer Science*, pages 297–313. Springer, 2015. doi: 10.1007/978-3-319-19069-3_19. URL https://doi.org/10.1007/978-3-319-19069-3_19.
- [9] Rada Mihalcea. *The Text Mining Handbook: Advanced Approaches to Analyzing Unstructured Data* ronon feldman and james sanger (bar-ilan university and ABS ventures) cambridge, england: Cambridge university press, 2007, xii+410 pp; hardbound, ISBN 0-521-83657-3. *Comput. Linguistics*, 34(1):125–127, 2008. doi: 10.1162/coli.2008.34.1.125. URL <https://doi.org/10.1162/coli.2008.34.1.125>.
- [10] Quoc V. Le and Tomas Mikolov. Distributed representations of sentences and documents. In *Proceedings of the 31th International Conference on Machine Learning, ICML 2014, Beijing, China, 21-26 June 2014*, volume 32 of *JMLR Workshop and Conference Proceedings*, pages 1188–1196. JMLR.org, 2014. URL <http://proceedings.mlr.press/v32/le14.html>.
- [11] David M. Blei, Andrew Y. Ng, and Michael I. Jordan. Latent dirichlet allocation. *J. Mach. Learn. Res.*, 3:993–1022, 2003. URL <http://jmlr.org/papers/v3/blei03a.html>.
- [12] Dan Jurafsky and James H. Martin. *Speech and language processing: an introduction to natural language processing, computational linguistics, and speech recognition, 2nd Edition*. Prentice Hall series in artificial intelligence. Prentice Hall, Pearson Education International, 2009. ISBN 9780135041963. URL <http://www.worldcat.org/oclc/315913020>.
- [13] Stuart J. Russell and Peter Norvig. *Artificial Intelligence - A Modern Approach, Third International Edition*. Pearson Education, 2010. ISBN

- 978-0-13-207148-2. URL http://vig.pearsoned.com/store/product/1,1207,store-12521_isbn-0136042597,00.html.
- [14] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural Computation*, 9(8):1735–1780, 1997. doi: 10.1162/neco.1997.9.8.1735. URL <https://doi.org/10.1162/neco.1997.9.8.1735>.
 - [15] Razvan Pascanu, Tomas Mikolov, and Yoshua Bengio. On the difficulty of training recurrent neural networks. In *Proceedings of the 30th International Conference on Machine Learning, ICML 2013, Atlanta, GA, USA, 16-21 June 2013*, volume 28 of *JMLR Workshop and Conference Proceedings*, pages 1310–1318. JMLR.org, 2013. URL <http://proceedings.mlr.press/v28/pascanu13.html>.
 - [16] Felix A. Gers, Jürgen Schmidhuber, and Fred A. Cummins. Learning to forget: Continual prediction with LSTM. *Neural Computation*, 12(10):2451–2471, 2000. doi: 10.1162/089976600300015015. URL <https://doi.org/10.1162/089976600300015015>.
 - [17] Klaus Greff, Rupesh Kumar Srivastava, Jan Koutník, Bas R. Steunebrink, and Jürgen Schmidhuber. LSTM: A search space odyssey. *IEEE Trans. Neural Networks Learn. Syst.*, 28(10):2222–2232, 2017. doi: 10.1109/TNNLS.2016.2582924. URL <https://doi.org/10.1109/TNNLS.2016.2582924>.
 - [18] Yonghui Wu, Mike Schuster, Zhifeng Chen, Quoc V. Le, Mohammad Norouzi, Wolfgang Macherey, Maxim Krikun, Yuan Cao, Qin Gao, Klaus Macherey, Jeff Klingner, Apurva Shah, Melvin Johnson, Xiaobing Liu, Lukasz Kaiser, Stephan Gouws, Yoshikiyo Kato, Taku Kudo, Hideto Kazawa, Keith Stevens, George Kurian, Nishant Patil, Wei Wang, Cliff Young, Jason Smith, Jason Riesa, Alex Rudnick, Oriol Vinyals, Greg Corrado, Macduff Hughes, and Jeffrey Dean. Google’s neural machine translation system: Bridging the gap between human and machine translation. *CoRR*, abs/1609.08144, 2016. URL <http://arxiv.org/abs/1609.08144>.
 - [19] Kyunghyun Cho, Bart van Merriënboer, Çaglar Gülçehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning phrase representations using RNN encoder-decoder for statistical machine translation. In Alessandro Moschitti, Bo Pang, and Walter Daelemans, editors, *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing, EMNLP 2014, October 25-29, 2014, Doha, Qatar, A meeting of SIGDAT, a Special Interest Group of the ACL*, pages 1724–1734. ACL, 2014. doi: 10.3115/v1/d14-1179. URL <https://doi.org/10.3115/v1/d14-1179>.
 - [20] Denny Britz, Anna Goldie, Minh-Thang Luong, and Quoc V. Le. Massive exploration of neural machine translation architectures. *CoRR*, abs/1703.03906, 2017. URL <http://arxiv.org/abs/1703.03906>.
 - [21] Jürgen Schmidhuber. Deep learning in neural networks: An overview. *Neural Networks*, 61:85–117, 2015. doi: 10.1016/j.neunet.2014.09.003. URL <https://doi.org/10.1016/j.neunet.2014.09.003>.
 - [22] David E Rumelhart, Geoffrey E Hinton, and Ronald J Williams. Learning representations by back-propagating errors. *nature*, 323(6088):533–536, 1986.
 - [23] Chris Olah. Understanding LSTM networks, 2015. URL <http://>

- [//web.archive.org/web/20200528192048/http://colah.github.io/posts/2015-08-Understanding-LSTMs/](http://web.archive.org/web/20200528192048/http://colah.github.io/posts/2015-08-Understanding-LSTMs/). Accessed = 2020-06-02.
- [24] Boudewijn F. van Dongen, R. A. Crooy, and Wil M. P. van der Aalst. Cycle time prediction: When will this case finally be finished? In Robert Meersman and Zahir Tari, editors, *On the Move to Meaningful Internet Systems: OTM 2008, OTM 2008 Confederated International Conferences, CoopIS, DOA, GADA, IS, and ODBASE 2008, Monterrey, Mexico, November 9-14, 2008, Proceedings, Part I*, volume 5331 of *Lecture Notes in Computer Science*, pages 319–336. Springer, 2008. doi: 10.1007/978-3-540-88871-0_22. URL https://doi.org/10.1007/978-3-540-88871-0_22.
 - [25] Wil M. P. van der Aalst, M. H. Schonenberg, and Minseok Song. Time prediction based on process mining. *Inf. Syst.*, 36(2):450–475, 2011. doi: 10.1016/j.is.2010.09.001. URL <https://doi.org/10.1016/j.is.2010.09.001>.
 - [26] Suraj Pandey, Surya Nepal, and Shipping Chen. A test-bed for the evaluation of business process prediction techniques. In Dimitrios Georgakopoulos and James B. D. Joshi, editors, *7th International Conference on Collaborative Computing: Networking, Applications and Worksharing, CollaborateCom 2011, Orlando, FL, USA, 15-18 October, 2011*, pages 382–391. ICST / IEEE, 2011. doi: 10.4108/icst.collaboratecom.2011.247129. URL <https://doi.org/10.4108/icst.collaboratecom.2011.247129>.
 - [27] Andreas Rogge-Solti and Mathias Weske. Prediction of remaining service execution time using stochastic petri nets with arbitrary firing delays. In Samik Basu, Cesare Pautasso, Liang Zhang, and Xiang Fu, editors, *Service-Oriented Computing - 11th International Conference, ICSOC 2013, Berlin, Germany, December 2-5, 2013, Proceedings*, volume 8274 of *Lecture Notes in Computer Science*, pages 389–403. Springer, 2013. doi: 10.1007/978-3-642-45005-1_27. URL https://doi.org/10.1007/978-3-642-45005-1_27.
 - [28] Michelangelo Ceci, Pasqua Fabiana Lanotte, Fabio Fumarola, Dario Pietro Cavallo, and Donato Malerba. Completion time and next activity prediction of processes using sequential pattern mining. In Saso Dzeroski, Pance Panov, Dragi Kocev, and Ljupco Todorovski, editors, *Discovery Science - 17th International Conference, DS 2014, Bled, Slovenia, October 8-10, 2014. Proceedings*, volume 8777 of *Lecture Notes in Computer Science*, pages 49–61. Springer, 2014. doi: 10.1007/978-3-319-11812-3_5. URL https://doi.org/10.1007/978-3-319-11812-3_5.
 - [29] Irene Teinemaa, Marlon Dumas, Fabrizio Maria Maggi, and Chiara Di Francesco-marino. Predictive business process monitoring with structured and unstructured data. In Marcello La Rosa, Peter Loos, and Oscar Pastor, editors, *Business Process Management - 14th International Conference, BPM 2016, Rio de Janeiro, Brazil, September 18-22, 2016. Proceedings*, volume 9850 of *Lecture Notes in Computer Science*, pages 401–417. Springer, 2016. doi: 10.1007/978-3-319-45348-4_23. URL https://doi.org/10.1007/978-3-319-45348-4_23.
 - [30] Joerg Evermann, Jana-Rebecca Rehse, and Peter Fettke. A deep learning approach for predicting process behaviour at runtime. In Marlon Dumas and Marcelo Fantinato, editors, *Business Process Management Workshops - BPM 2016 International Workshops, Rio de Janeiro, Brazil, September 19, 2016, Revised Papers*, volume 281 of

- Lecture Notes in Business Information Processing*, pages 327–338, 2016. doi: 10.1007/978-3-319-58457-7_24. URL https://doi.org/10.1007/978-3-319-58457-7_24.
- [31] Niek Tax, Ilya Verenich, Marcello La Rosa, and Marlon Dumas. Predictive business process monitoring with LSTM neural networks. In Eric Dubois and Klaus Pohl, editors, *Advanced Information Systems Engineering - 29th International Conference, CAiSE 2017, Essen, Germany, June 12-16, 2017, Proceedings*, volume 10253 of *Lecture Notes in Computer Science*, pages 477–492. Springer, 2017. doi: 10.1007/978-3-319-59536-8_30. URL https://doi.org/10.1007/978-3-319-59536-8_30.
- [32] Nicolò Navarin, Beatrice Vincenzi, Mirko Polato, and Alessandro Sperduti. LSTM networks for data-aware remaining time prediction of business process instances. In *2017 IEEE Symposium Series on Computational Intelligence, SSCI 2017, Honolulu, HI, USA, November 27 - Dec. 1, 2017*, pages 1–7. IEEE, 2017. doi: 10.1109/SSCI.2017.8285184. URL <https://doi.org/10.1109/SSCI.2017.8285184>.
- [33] Mirko Polato, Alessandro Sperduti, Andrea Burattin, and Massimiliano de Leoni. Time and activity sequence prediction of business process instances. *Computing*, 100(9):1005–1031, 2018. doi: 10.1007/s00607-018-0593-x. URL <https://doi.org/10.1007/s00607-018-0593-x>.
- [34] Irene Teinemaa, Marlon Dumas, Marcello La Rosa, and Fabrizio Maria Maggi. Outcome-oriented predictive process monitoring: Review and benchmark. *ACM Trans. Knowl. Discov. Data*, 13(2):17:1–17:57, 2019. doi: 10.1145/3301300. URL <https://doi.org/10.1145/3301300>.
- [35] Gyunam Park and Minseok Song. Prediction-based resource allocation using LSTM and minimum cost and maximum flow algorithm. In *International Conference on Process Mining, ICPM 2019, Aachen, Germany, June 24-26, 2019*, pages 121–128. IEEE, 2019. doi: 10.1109/ICPM.2019.00027. URL <https://doi.org/10.1109/ICPM.2019.00027>.
- [36] R. P. Jagadeesh Chandra Bose, Wil M. P. van der Aalst, Indre Zliobaite, and Mykola Pechenizkiy. Dealing with concept drifts in process mining. *IEEE Trans. Neural Networks Learn. Syst.*, 25(1):154–171, 2014. doi: 10.1109/TNNLS.2013.2278313. URL <https://doi.org/10.1109/TNNLS.2013.2278313>.
- [37] Zellig S Harris. Distributional structure. *Word*, 10(2-3):146–162, 1954.
- [38] Yoshua Bengio, Réjean Ducharme, Pascal Vincent, and Christian Janvin. A neural probabilistic language model. *J. Mach. Learn. Res.*, 3:1137–1155, 2003. URL <http://jmlr.org/papers/v3/bengio03a.html>.
- [39] Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient estimation of word representations in vector space. In Yoshua Bengio and Yann LeCun, editors, *1st International Conference on Learning Representations, ICLR 2013, Scottsdale, Arizona, USA, May 2-4, 2013, Workshop Track Proceedings*, 2013. URL <http://arxiv.org/abs/1301.3781>.
- [40] Lei Jimmy Ba, Jamie Ryan Kiros, and Geoffrey E. Hinton. Layer normalization. *CoRR*, abs/1607.06450, 2016. URL <http://arxiv.org/abs/1607.06450>.

- [41] Mahidhar Dwarampudi and N. V. Subba Reddy. Effects of padding on lstms and cnns. *CoRR*, abs/1903.07288, 2019. URL <http://arxiv.org/abs/1903.07288>.
- [42] Timothy Dozat. Incorporating nesterov momentum into adam. 2016.

Acknowledgments